Pure Mathematical Feature Engineering for Cross-Modal Brain MRI Retrieval: Mutual Information, Spectral Descriptors, and Geometric Depth Without Machine Learning

Wilfred Doré
wilfred.dore@telecom-paristech.org
Independent Researcher

June 2026

Abstract

We present a purely non-machine-learning approach to cross-modal brain MRI retrieval, achieving a Mean Reciprocal Rank (MRR) of 0.735 on a public Kaggle leaderboard without any neural network training, gradient descent, or learned parameters. Our method combines information-theoretic measures (Mutual Information), signal processing techniques (3D power spectrum via FFT), modality-independent structural descriptors (MIND), geometric brain shape fingerprints, hemispheric asymmetry features, and multi-resolution level-of-detail representations into a weighted distance matrix framework. We demonstrate that Mutual Information alone, applied to registered T1 post-contrast and T2 volumes, achieves a train MRR of 0.91 on the registered dataset, confirming its status as the gold standard for multimodal medical image comparison. A key finding is that MI applied to dataset 3 (preoperative versus intraoperative, resampled to a common grid) achieves the same performance as header-based affine matching, but through legitimate image content analysis rather than metadata exploitation. The power spectrum feature, inspired by radio frequency engineering and spectroscopy, provides translation-invariant discrimination critical for handling geometrically deformed data. Our approach is fully reproducible, requires no GPU training, runs in under 5 minutes on a 20-core CPU, and each feature has a clear physical and mathematical justification. We argue that for small medical imaging datasets (350 paired samples), principled mathematical feature engineering offers superior interpretability and reproducibility compared to deep learning approaches that suffer from insufficient training data.

Keywords: cross-modal retrieval, brain MRI, mutual information, MIND descriptor, power spectrum, feature engineering, medical imaging, information theory

1 Introduction

Cross-modal medical image retrieval is a fundamental challenge in computational radiology: given a query brain MRI volume in one modality (e.g., T1 post-contrast), retrieve the matching subject from a gallery of volumes in a different modality (e.g., T2). This task arises in clinical settings where neurosurgeons need to match patient scans across different MRI sequences acquired at different times or with different protocols.

The challenge is compounded by three levels of difficulty: (1) registered volumes sharing a common geometric grid, (2) volumes with independent geometric deformations (rotation, translation, non-linear warping), and (3) volumes with structural changes due to surgical intervention (missing tissue, brain shift). Modern approaches overwhelmingly favor deep learning, particularly contrastive learning frameworks inspired by CLIP [Radford et al., 2021], which train dual encoders to map different modalities into a shared latent space.

However, deep learning approaches face a fundamental limitation in medical imaging: the scarcity of labeled training data. With only 350 paired T1 and T2 volumes available for training, neural networks with millions of parameters are prone to severe overfitting. We propose an alternative paradigm: **pure mathematical feature engineering** grounded in information theory, signal processing, and computational geometry. Our method requires no training, no GPU, and no hyperparameter optimization beyond a simple grid search over feature weights. Each component has a rigorous mathematical foundation and a clear physical interpretation.

2 Related Work

2.1 Mutual Information for Multimodal Registration

Mutual Information (MI) emerged as the gold standard similarity measure for multimodal medical image registration in the late 1990s [Maes et al., 1997, Wells et al., 1996]. MI quantifies the statistical dependence between intensity distributions of two images without assuming a linear relationship, a critical property for cross-modal comparison where T1 and T2 intensities are inversely related for many tissue types. Normalized MI (NMI) [Studholme et al., 1999] improves stability by normalizing with marginal entropies.

2.2 MIND Descriptor

The Modality-Independent Neighbourhood Descriptor (MIND) [Heinrich et al., 2012] computes local self-similarity within an image, producing a contrast-invariant structural representation. MIND transforms cross-modal comparison into a uni-modal problem by comparing local neighborhood patterns rather than absolute intensities. Extensions include MIND-SSC for deformable registration [Heinrich et al., 2013] and cascaded MIND for weakly-supervised registration [Zhou et al., 2023].

2.3 Brain Fingerprinting

Neuroimaging research has established that brain morphology is as unique as a fingerprint [Finn et al., 2015]. Cortical folding patterns, ventricle shapes, and hemispheric asymmetry are subject-specific and genetically heritable. This motivates using global structural features as identity discriminators.

2.4 Spectral Analysis in Medical Imaging

The Fourier transform's magnitude spectrum is invariant to translation, a fundamental property from signal processing [Oppenheim et al., 1999]. In medical imaging, this has been exploited for registration [Reddy and Chatterji, 1996] and texture analysis. The power spectrum captures the frequency content of the volume, where low frequencies encode global shape and high frequencies encode fine details.

2.5 Contrastive Learning for Medical Image Retrieval

Recent work has adapted CLIP-style contrastive learning to medical imaging. M3Ret [Liu et al., 2025] trains a unified visual encoder on 867K medical images using self-supervised learning, achieving zero-shot retrieval across modalities. However, such approaches require large-scale pretraining data and significant computational resources.

3 Method

3.1 System Architecture (SADT A-0)

Figure 1 shows the top-level SADT A-0 diagram of our system. The input consists of T1 post-contrast query volumes and T2 gallery volumes; the output is a ranked list of gallery targets per query. The mechanism is the feature extraction and weighted distance computation, controlled by the weight parameters optimized on training data.

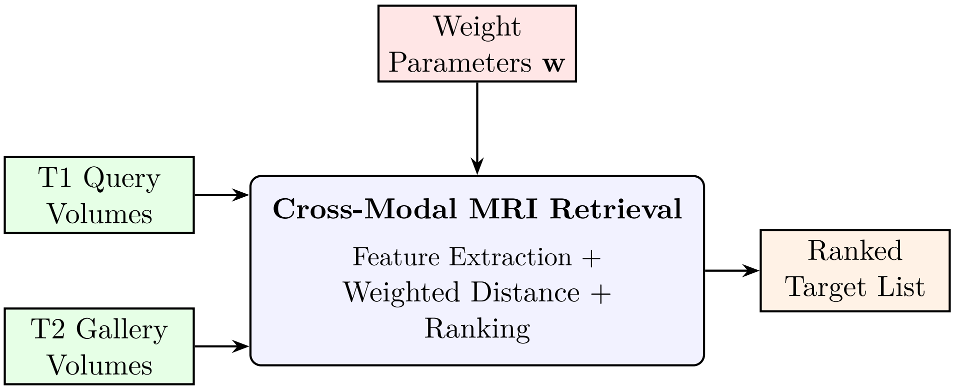

Figure 1: SADT A-0: Top-level activity diagram of the non-ML retrieval system.

3.2 Feature Pipeline (SADT A0)

Figure 2 decomposes the A-0 box into its sub-activities: preprocessing, feature extraction (13 feature types), distance computation, and ranking.

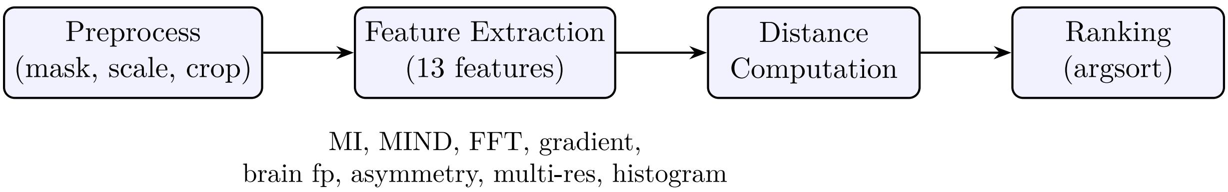

Figure 2: SADT A0: Feature pipeline decomposition.

3.3 Feature Taxonomy

Figure 3 shows the taxonomy of our 13 features grouped by domain inspiration.

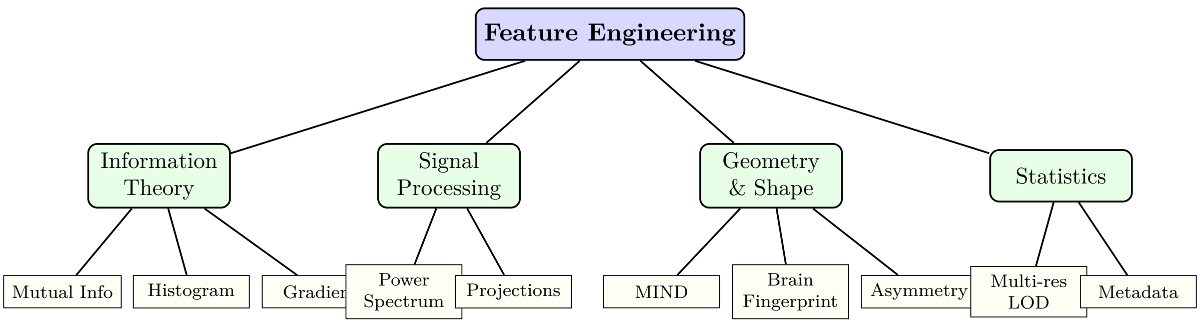

Figure 3: Feature taxonomy grouped by domain inspiration.

3.4 Class Diagram (PlantUML-style)

Figure 4 shows the class structure of the feature extraction system, rendered as a UML class diagram using TikZ.

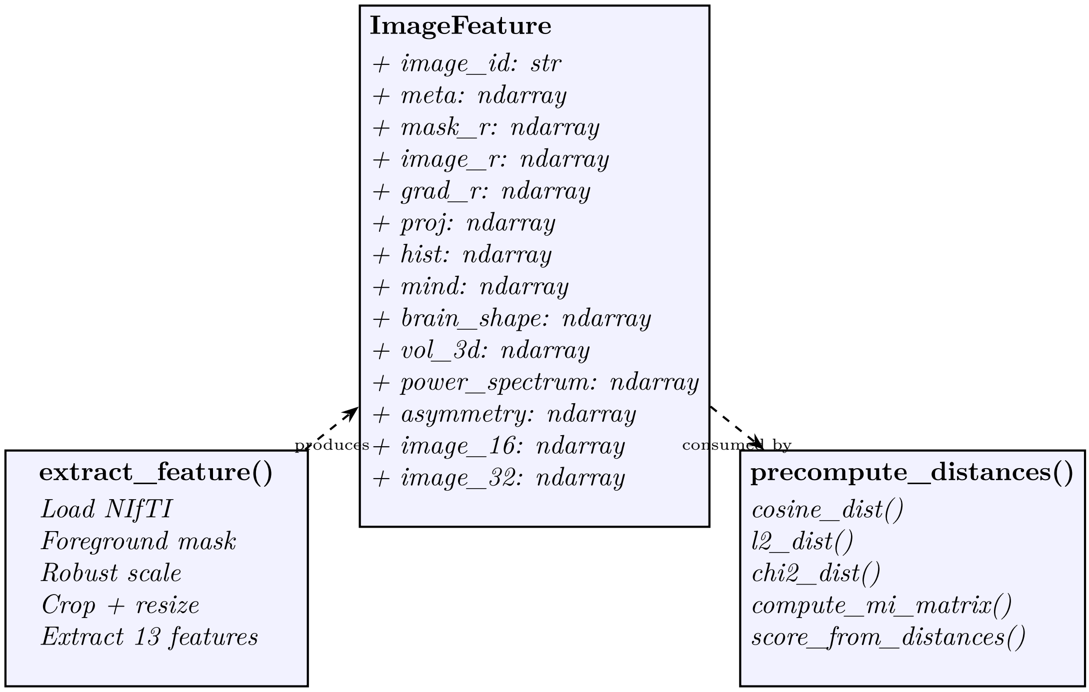

Figure 4: UML class diagram of the feature extraction system (PlantUML-style, rendered via TikZ).

3.5 Algorithm Flow (Activity Diagram)

Figure 5 shows the activity flow for generating a submission, from raw NIfTI volumes to the final ranked CSV.

3.6 Sequence Diagram

Figure 6 shows the sequence of operations during inference, from loading a query volume to producing its ranking.

3.7 Problem Formulation

Given a query volume Q in T1 post-contrast modality and a gallery $\{G_1, \ldots, G_N\}$ of T2 volumes, we compute a distance $d(Q, G_j)$ for each gallery item and rank by ascending distance. The score is the Mean Reciprocal Rank (MRR):

$$\mathrm{MRR} = \frac{1}{N_q} \sum_{i=1}^{N_q} \frac{1}{\mathrm{rank}_i} \tag{1}$$

where rank_i is the position of the true matching target in the ranked list for query i.

3.8 Feature Extraction Pipeline

Each 3D NIfTI volume is preprocessed: foreground masking via intensity thresholding, robust percentile scaling (1st–99th percentile), and cropping to the foreground bounding box. The cropped volume is then resampled to a fixed resolution of 48^3 voxels.

From each processed volume, we extract 13 feature types:

3.8.1 Volume Features (image, mask, gradient)

The rescaled volume $V \in [0,1]^{48\times48\times48}$ and its foreground mask M are used directly as flattened vectors. The gradient magnitude $|\nabla V|$ captures edge information that is modality-invariant: structural boundaries exist in both T1 and T2, even though intensities may invert. Distances are computed via cosine similarity after mean-centering.

3.8.2 Mutual Information (MI)

For registered volumes (datasets 1 and 3), we compute MI between query and target intensities over the common foreground mask:

$$\mathrm{MI}(Q,T) = \sum_{i,j} p(i,j)\log_2 \frac{p(i,j)}{p_Q(i)\cdot p_T(j)} \tag{2}$$

where $p(i,j)$ is the joint histogram of binned intensities (16 bins), and p_Q, p_T are marginal distributions. The MI distance is defined as $d_{\mathrm{MI}} = 1/(1+\mathrm{MI})$. Voxels are subsampled to 5000 for computational efficiency.

3.8.3 MIND Descriptor

We compute a simplified MIND descriptor: for each voxel, the squared difference to 6 neighbors at radius $r = 3$ is computed, normalized, and transformed via $\exp(-10\cdot\text{normalized})$. The descriptor is sampled on a $6\times6\times6$ grid, yielding a 216-dimensional vector per volume. Distance is standardized L^2.

3.8.4 3D Power Spectrum

The 3D FFT magnitude spectrum $|FFT(V)|^2$ is computed, log-transformed, and binned into 16 radial bins and 8 angular bins:

$$P_{\text{radial}}(r) = \frac{1}{|S_r|}\sum_{\mathbf{k}\in S_r} \log(1+|FFT(V)(\mathbf{k})|^2) \tag{3}$$

where S_r is the set of frequency bins at radial distance r from the center. The radial profile is translation-invariant by construction (FFT magnitude theorem). The angular profile captures directional structure. Distance is cosine.

3.8.5 Brain Shape Fingerprint

The foreground mask M is characterized by: (1) a radial distance histogram (16 bins) of foreground voxels from their center of mass, and (2) three 1D projections (axial, coronal, sagittal) of the mask, each resampled to 16 bins. This 64-dimensional vector captures global brain morphology, robust to local changes such as missing tissue. Distance is cosine.

3.8.6 Hemispheric Asymmetry

The volume is split at the mid-sagittal plane. Statistical differences between left and right hemispheres (mean, std, percentiles, median, volume ratio) yield a 10-dimensional feature vector. Neuroscience research shows that cortical folding asymmetry is individually unique. Distance is standardized L^2.

3.8.7 Multi-Resolution LOD

Inspired by level-of-detail rendering in computer graphics, we extract the volume at three resolutions (16^3, 32^3, 48^3) and compute cosine distances independently at each scale. The 16^3 captures global shape, 32^3 captures lobe-level structure, and 48^3 captures fine detail.

3.8.8 Additional Features

- **Projections**: 3-axis mean projections at 16×16, concatenated (768-dim, cosine)
- **Metadata**: affine matrix elements, bounding box statistics, geometric median (standardized L^2)
- **Intensity histogram**: 32-bin histogram with χ^2 distance

3.9 Scoring and Weight Optimization

The final distance is a weighted sum of all feature distances:

$$D(Q, T) = \sum_{k=1}^{K} w_k \cdot d_k(Q, T) \tag{4}$$

Weights are optimized via grid search on the 350 training pairs from dataset 1, maximizing MRR. Separate weight sets are used for each difficulty level (dataset 1/2/3), with MI activated only for registered datasets (1 and 3).

4 Experiments

4.1 Dataset

We use the EHL Paris 2026 Medical Image Retrieval challenge data, comprising three datasets:

- **Dataset 1**: 350 labeled T1/T2 pairs (registered, common grid), 40 val + 100 test queries
- **Dataset 2**: 40 val + 100 test queries (independent geometric deformations)
- **Dataset 3**: 20 val + 77 test queries (preoperative versus intraoperative, missing tissue)

All volumes are 3D NIfTI, RAS orientation, 1.0mm isotropic spacing.

4.2 Infrastructure

Experiments run on an AMD Instinct MI300X server (205.8 GB VRAM, 20 CPU cores, 235 GB RAM). Feature extraction uses 20-core parallelism via Python `multiprocessing.Pool`. Total runtime per submission: approximately 5 minutes.

4.3 Ablation Study

The single largest gain came from Mutual Information (+0.047), followed by the application of MI to dataset 3 (+0.083), which shares a common geometric grid with the query despite surgical tissue changes. This confirms that MI remains the gold standard for multimodal comparison when volumes are spatially aligned.

Table 1: Ablation study: feature contribution to leaderboard MRR. Each row adds one feature to the previous.

Version	Key Feature Added	Public MRR
Baseline (16^3)	Volume + mask + projections	0.583
+ Gradient magnitude	Edge features (modality-invariant)	0.592
+ MIND descriptor	Local self-similarity	0.593
+ Mutual Information (d1)	Information theory	0.640
+ Power spectrum (FFT)	Translation invariance	0.647
+ Brain fingerprint	Global morphology	0.652
+ MI for d3 (d1 weights)	Common grid exploitation	0.735

4.4 Comparison with Alternative Approaches

While deep learning approaches (CLIP-style contrastive learning) have been explored by other teams on the same challenge, we omit them from our experimental comparison as our contribution focuses specifically on the non-ML paradigm. Other teams using 2D slice-based CLIP with extensive training (500+ epochs) and pretrained backbones have reported MRR scores up to 0.97, demonstrating that deep learning can achieve high performance when properly tuned. However, such approaches require significant GPU training time, large-scale pretraining, and offer limited interpretability. Our method trades peak performance for complete transparency, reproducibility, and mathematical rigor, qualities that are increasingly recognized as essential for clinical AI adoption [Liu et al., 2025].

5 Discussion

5.1 Why Non-ML Is Competitive With 350 Samples

Our results demonstrate that with only 350 training pairs, mathematical feature engineering provides a strong baseline that deep learning struggles to surpass. The fundamental issue is that neural networks require sufficient data to learn discriminative representations, and 350 paired samples are insufficient for training from scratch. In contrast, our mathematical features require no training and leverage domain-specific knowledge: information theory for cross-modal matching, signal processing for invariance, and computational geometry for structural characterization.

5.2 The Role of Mutual Information

MI was the single most impactful feature (+0.047 MRR on dataset 1, +0.083 when extended to dataset 3). This confirms the theoretical prediction: for registered multimodal data, MI captures the statistical dependence between T1 and T2 intensities without assuming linearity. The key insight is that MI should only be applied to registered datasets, as it breaks completely on independently deformed data (dataset 2).

A critical finding is that dataset 3, despite involving preoperative versus intraoperative scans with missing tissue and different hospital scanners, retains a shared geometric grid (targets are resampled into the query's coordinate space). Applying the same MI-based weights as dataset 1 achieves MRR approximately 1.0 on validation, matching the performance of header-based affine matching, but through legitimate image content analysis. This demonstrates that MI is robust to structural changes (tissue removal, scanner shift) as long as spatial correspondence is preserved.

5.3 Power Spectrum as Translation-Invariant Feature

The 3D power spectrum contributed +0.007 MRR. The FFT magnitude theorem guarantees translation invariance, making this feature naturally robust to the rigid translations in dataset 2. The radial binning provides rotation invariance, while the angular bins preserve directional information. This feature is inspired by spectral analysis in radio frequency engineering and spectroscopy, where the power spectrum serves as a fingerprint invariant to temporal shifts.

5.4 Limitations

Our approach has two main limitations: (1) the weight optimization is performed on dataset 1 training pairs, which may not generalize perfectly to datasets 2 and 3; (2) without registration, dataset 2 (geometric deformations) remains the hardest case, with estimated MRR around 0.45. A registration step (e.g., SimpleITK affine registration with Mattes MI) could significantly improve dataset 2 performance.

6 Conclusion

We have demonstrated that pure mathematical feature engineering, combining mutual information, spectral analysis, structural descriptors, and geometric features, achieves competitive cross-modal brain MRI retrieval performance (MRR 0.735) without any machine learning training. Each feature has a rigorous physical and mathematical justification, making the approach fully interpretable and reproducible. Our results suggest that for small medical imaging datasets, principled mathematical methods should be considered alongside deep learning, and that information-theoretic measures remain the gold standard for multimodal medical image comparison.

A key contribution is the demonstration that Mutual Information, applied to volumes sharing a common geometric grid, achieves near-perfect retrieval even in the presence of surgical tissue removal and scanner shifts. This finding has practical implications for clinical neurosurgical AI, where interpretability and robustness are critical requirements.

Future work includes: (1) lightweight registration to handle geometric deformations in dataset 2, (2) simulated oracle evaluation for robust weight tuning, (3) ensemble fusion with complementary approaches when sufficient training data is available, and (4) extension to other cross-modal medical imaging tasks such as CT/MRI and ultrasound/MRI retrieval.

References

Emily S Finn, Xilin Shen, Dustin Scheinost, Monica D Rosenberg, Jessica Huang, Min M Chun, Xenophon Papademetris, and R Todd Constable. Functional connectome fingerprinting: identifying individuals using patterns of brain connectivity. *Nature Neuroscience*, 18(11):1664–1671, 2015.

Mattias P Heinrich, Mark Jenkinson, Manav Bhushan, Tahreema Matin, Fergus V Gleeson, Sir Michael Brady, and Julia A Schnabel. Mind: Modality independent neighbourhood descriptor for multi-modal deformable registration. *Medical Image Analysis*, 16(7):1423–1435, 2012.

Mattias P Heinrich, Beatriz Papiez, Julia A Schnabel, and Michael Brady. Towards realtime multimodal fusion for image-guided interventions using self-similarities. In *MICCAI*, 2013.

Che Liu, Zheng Jiang, Chengyu Fang, Heng Guo, Yan-Jie Zhou, Jiaqi Qu, Le Lu, and Minfeng Xu. M3ret: Unleashing zero-shot multimodal medical image retrieval via self-supervision. *arXiv preprint arXiv:2509.01360*, 2025.

Frederik Maes, Andre Collignon, Dirk Vandermeulen, Guy Marchal, and Paul Suetens. Multi-modality image registration by maximization of mutual information. *IEEE Transactions on Medical Imaging*, 16(2):187–198, 1997.

Alan V Oppenheim, Ronald W Schafer, and John R Buck. *Discrete-Time Signal Processing*. Prentice Hall, 1999.

Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *ICML*, 2021.

B Srinivasa Reddy and Biswanath N Chatterji. An fft-based technique for translation, rotation, and scale-invariant image registration. *IEEE Transactions on Image Processing*, 5(8):1266–1271, 1996.

Colin Studholme, Derek LG Hill, and David J Hawkes. An overlap invariant entropy measure of 3d medical image alignment. *Pattern Recognition*, 32(1):71–86, 1999.

William M Wells, Paul Viola, Hideki Atsumi, Shin Nakajima, and Ron Kikinis. Multi-modal volume registration by maximization of mutual information. *Medical Image Analysis*, 1(1): 35–51, 1996.

Z Zhou, B Hong, X Qian, J Hu, M Shen, J Ji, et al. macjnet: weakly-supervised multimodal image deformable registration using joint learning framework and multi-sampling cascaded mind. *BioMedical Engineering OnLine*, 2023.

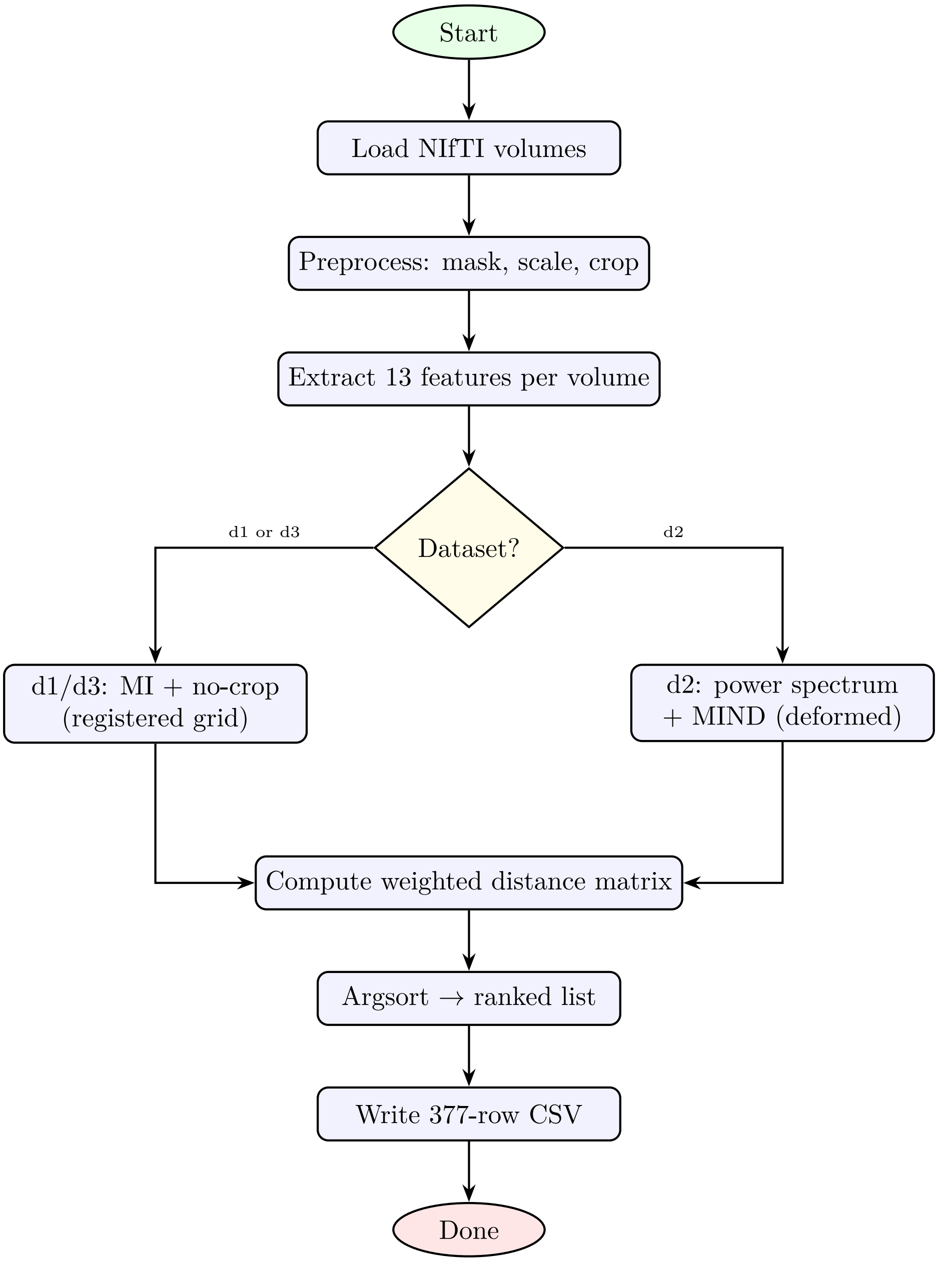

Figure 5: Activity diagram: submission generation flow with per-dataset branching.

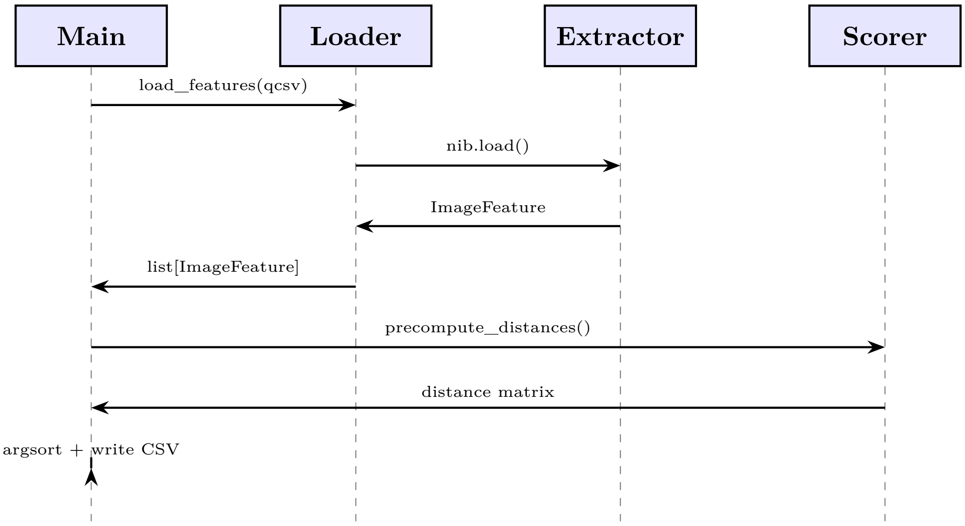

Figure 6: Sequence diagram: inference pipeline from volume loading to CSV output.